

```
In [ ]: %pip install matplotlib numpy pandas
```

```
Requirement already satisfied: matplotlib in e:\practice\pattern_2024\venv\lib\site-packages (3.8.2)
Requirement already satisfied: numpy in e:\practice\pattern_2024\venv\lib\site-packages (1.26.3)
Requirement already satisfied: pandas in e:\practice\pattern_2024\venv\lib\site-packages (2.1.4)
Requirement already satisfied: kiwisolver>=1.3.1 in e:\practice\pattern_2024\venv\lib\site-packages (from matplotlib) (1.4.5)
Requirement already satisfied: python-dateutil>=2.7 in e:\practice\pattern_2024\venv\lib\site-packages (from matplotlib) (2.8.2)
Requirement already satisfied: pyparsing>=2.3.1 in e:\practice\pattern_2024\venv\lib\site-packages (from matplotlib) (3.1.1)
Requirement already satisfied: fonttools>=4.22.0 in e:\practice\pattern_2024\venv\lib\site-packages (from matplotlib) (4.47.2)
Requirement already satisfied: packaging>=20.0 in e:\practice\pattern_2024\venv\lib\site-packages (from matplotlib) (23.2)
Requirement already satisfied: cyclor>=0.10 in e:\practice\pattern_2024\venv\lib\site-packages (from matplotlib) (0.12.1)
Requirement already satisfied: pillow>=8 in e:\practice\pattern_2024\venv\lib\site-packages (from matplotlib) (10.2.0)
Requirement already satisfied: contourpy>=1.0.1 in e:\practice\pattern_2024\venv\lib\site-packages (from matplotlib) (1.2.0)
Requirement already satisfied: pytz>=2020.1 in e:\practice\pattern_2024\venv\lib\site-packages (from pandas) (2023.3.post1)
Requirement already satisfied: tzdata>=2022.1 in e:\practice\pattern_2024\venv\lib\site-packages (from pandas) (2023.4)
Requirement already satisfied: six>=1.5 in e:\practice\pattern_2024\venv\lib\site-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
Note: you may need to restart the kernel to use updated packages.
```

```
[notice] A new release of pip is available: 23.0.1 -> 23.3.2
```

```
[notice] To update, run: python.exe -m pip install --upgrade pip
```

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [ ]: train_url = "http://s3.amazonaws.com/assets.datacamp.com/course/Kaggle/train.csv"
train = pd.read_csv(train_url) # training set

test_url = "http://s3.amazonaws.com/assets.datacamp.com/course/Kaggle/test.csv"
test = pd.read_csv(test_url) # test set
```

```
In [ ]: print(train.head())
print(train.tail())
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

	PassengerId	Survived	Pclass	Name	\
886	887	0	2	Montvila, Rev. Juozas	
887	888	1	1	Graham, Miss. Margaret Edith	
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	
889	890	1	1	Behr, Mr. Karl Howell	
890	891	0	3	Dooley, Mr. Patrick	

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
886	male	27.0	0	0	211536	13.00	NaN	S
887	female	19.0	0	0	112053	30.00	B42	S
888	female	NaN	1	2	W./C. 6607	23.45	NaN	S
889	male	26.0	0	0	111369	30.00	C148	C
890	male	32.0	0	0	370376	7.75	NaN	Q

```
In [ ]: train.describe()
```

```
Out[ ]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

T8

```
In [ ]: # T8
median_age = train["Age"].median()

print(f"Median of age is {median_age}")

train["Age"] = train["Age"].fillna(train["Age"].median())

Median of age is 28.0
```

T9

```
In [ ]: # T9
train["Embarked"] = train["Embarked"].fillna(train["Embarked"].mode()[0])

train.loc[train["Embarked"] == "S", "Embarked"] = 0
train.loc[train["Embarked"] == "C", "Embarked"] = 1
train.loc[train["Embarked"] == "Q", "Embarked"] = 2

train.loc[train["Sex"] == "male", "Sex"] = 0
train.loc[train["Sex"] == "female", "Sex"] = 1
```

T10

```
In [ ]: features = np.array(train[["Pclass", "Sex", "Age", "Embarked"]].values, dtype=float)
results = np.array(train["Survived"].values, dtype=float)
features, results
```



```

def get_next_theta(theta, x, y):
    error = y - calculate_logistic(np.dot(x, theta))

    diff_theta = x.T.dot(error) * learning_rate

    return diff_theta

def classifier(theta, x):
    solution = calculate_logistic(np.dot(x, theta))

    classifier_result = solution >= 0.5

    solution[classifier_result == True] = 1
    solution[classifier_result == False] = 0

    return solution

def measurement(predicts, actual):
    diff = predicts - actual
    correct = diff == 0

    return np.sum(correct) / actual.shape

def train_logistic_regression(features, results):
    iterations = int(1e4)
    features_with_one = np.insert(features, 0, 1, axis=1)

    starting_theta = np.zeros(features_with_one.shape[1])
    theta = np.copy(starting_theta)
    accuracy_list = []

    for _ in range(iterations):
        theta += get_next_theta(theta, features_with_one, results)
        accuracy = measurement(
            np.array(train["Survived"]), classifier(theta, features_with_one)
        )[0]
        accuracy_list.append(accuracy)

    print(theta)
    print("Training Accuracy:", accuracy_list[-1])

    plt.plot(accuracy_list)
    plt.show()

    return theta


```

```

In [ ]: theta = train_logistic_regression(features, results)

[ 0.60877697 -0.77117063  2.1963073 -0.01308992  0.34021432]
Training Accuracy: 0.7946127946127947

```



```
In [ ]: # Clean test data
test["Embarked"] = test["Embarked"].fillna(test["Embarked"].mode()[0])

test.loc[test["Embarked"] == "S", "Embarked"] = 0
test.loc[test["Embarked"] == "C", "Embarked"] = 1
test.loc[test["Embarked"] == "Q", "Embarked"] = 2

test.loc[test["Sex"] == "male", "Sex"] = 0
test.loc[test["Sex"] == "female", "Sex"] = 1
```

T11

```
In [ ]: # T11
test_features = np.array(test[["Pclass", "Sex", "Age", "Embarked"]].values, dtype=float)
test_features = np.insert(test_features, 0, 1, axis=1)

test_result = classifier(theta, test_features)
test_id = test[["PassengerId"]]

df = pd.DataFrame()

df["PassengerId"] = test_id
df["Survived"] = np.array(test_result, dtype=int)

df.to_csv("titanic_result.csv", index=False)
```

Submission Details



titanic_result.csv

Complete · 2m ago

Score: 0.66985

UPLOADED FILES



titanic_result.csv (3 KiB)



DESCRIPTION

Logistic Regression with Gradient Descend V3


44 / 500

T12

```
In [ ]: features_with_higher_order = np.insert(features, 0, features[:, 0] ** 2, axis=1)
features_with_higher_order = np.insert(
    features_with_higher_order, 0, features[:, 0] * features[:, 2], axis=1
)

theta_higher_order = train_logistic_regression(features_with_higher_order, results)

[ 0.27778632 -0.01063309 -0.23897882  0.22627652  2.13072319 -0.00735009
  0.36545214]
Training Accuracy: 0.8058361391694725
```



```
In [ ]: test_features = np.array(test[["Pclass", "Sex", "Age", "Embarked"]].values, dtype=float)

test_features_higher_order = np.insert(
    test_features, 0, test_features[:, 0] ** 2, axis=1
)
```

```

test_features_higher_order = np.insert(
    test_features_higher_order, 0, test_features[:, 0] * test_features[:, 2], axis=1
)
test_features_higher_order = np.insert(test_features_higher_order, 0, 1, axis=1)

test_result = classifier(theta_higher_order, test_features_higher_order)
test_id = test[["PassengerId"]]


df = pd.DataFrame()

df["PassengerId"] = test_id
df["Survived"] = np.array(test_result, dtype=int)

df.to_csv("titanic_result_higher_order.csv", index=False)


```


X
Submission Details


titanic_result_higher_order.csv
Complete · 37m ago

Score: 0.77511

UPLOADED FILES


titanic_result_higher_order.csv (3 KiB)



DESCRIPTION

Logistic regression with gradient descent V4 with higher order

62 / 500

We see that there are little accuracy difference which higher order feature gives more accuracy than normal feature

T13


```

In [ ]: reduced_features = np.array(train[["Sex", "Age"]].values, dtype=float)

theta_reduced = train_logistic_regression(reduced_features, results)

[-0.68387827  2.04999267 -0.0178297 ]
Training Accuracy: 0.7867564534231201

```

```
In [ ]: test_features_reduced = np.array(test[["Sex", "Age"]].values, dtype=float)
test_features_reduced = np.insert(test_features_reduced, 0, 1, axis=1)

test_result = classifier(theta_reduced, test_features_reduced)
test_id = test[["PassengerId"]]

df = pd.DataFrame()

df["PassengerId"] = test_id
df["Survived"] = np.array(test_result, dtype=int)

df.to_csv("titanic_result_reduced.csv", index=False)
```

Submission Details



titanic_result_reduced.csv
Complete · 28m ago

Score: 0.77272

UPLOADED FILES

 titanic_result_reduced.csv (3 KiB)



DESCRIPTION

Logistic regression with gradient descent V4 + Reduced features to age, sex

75 / 500

Overall result from reduced features give less accuracy than normal feature or features which has higher order

OT3

```
In [ ]: # OT3
learning_rate = 1e-6
```

```

def measurement(results, actual):
    size = results.shape[0]
    return np.sum((results - actual) ** 2) / size

def get_next_theta(theta, x, y):
    diff_theta = x.T.dot(y - x.dot(theta)) * learning_rate

    return diff_theta

def train_linear_regression(features, results):
    iterations = int(1e6)
    features_with_one = np.insert(features, 0, 1, axis=1)

    starting_theta = np.zeros(features_with_one.shape[1])
    theta = np.copy(starting_theta)

    accuracy_list = []

    for _ in range(iterations):
        diff_theta = get_next_theta(theta, features_with_one, results)
        theta += diff_theta
        accuracy = measurement(
            np.array(train["Survived"]), np.dot(features_with_one, theta)
        )
        accuracy_list.append(accuracy)

    plt.plot(accuracy_list)
    plt.show()

    print("Mean Square Error:", accuracy_list[-1])
    return theta


```

```

In [ ]: linear_theta = train_linear_regression(features, results)
linear_theta

```

Mean Square Error: 0.1449257376613481



```

Out[ ]: array([ 0.77654442, -0.18843944,  0.49086711, -0.00505436,  0.04911346])

```

Result

- Mean Square Error: 0.1449257376613481
- Parameter: [0.77654442, -0.18843944, 0.49086711, -0.00505436, 0.04911346]

OT4

In []: # OT4

```
features_copy = np.insert(features, 0, 1, axis=1)

inverse_props = np.linalg.inv(np.matmul(features_copy.transpose(), features_copy))
linear_reg_theta = np.matmul(
    np.matmul(inverse_props, features_copy.transpose()), results
)

loss = measurement(np.array(train["Survived"]), np.dot(features_copy, linear_reg_theta))

print("Mean Square Error:", loss)
print("Parameter:", linear_reg_theta)
```

```
Mean Square Error: 0.14492573766134811
Parameter: [ 0.77654442 -0.18843944  0.49086711 -0.00505436  0.04911346]
```

Result

- Mean Square Error: 0.14492573766134811
- Parameter: [0.77654442 -0.18843944 0.49086711 -0.00505436 0.04911346]

which gives the same MSE and parameter value as in OT3