

Pattern_HW2_student_2024

February 3, 2024

0.1 For interactive notebook please see on [this link](#)

0.2 T1

$$\begin{aligned}
 T1: P(y_2, y_1, y_0 | \alpha) &= \frac{P(y_2, y_1, y_0, \alpha)}{P(\alpha)} \\
 &= \frac{P(y_2 | y_1, y_0, \alpha)}{P(\alpha)} P(y_1, y_0, \alpha) \\
 &= \frac{P(y_2 | y_1, y_0, \alpha)}{P(\alpha)} P(y_1 | y_0, \alpha) P(y_0, \alpha) \\
 &= P(y_2 | y_1, y_0, \alpha) P(y_1 | y_0, \alpha) P(y_0 | \alpha) \\
 P(y_2, y_1, y_0 | \alpha) &= \left(\frac{1}{\sqrt{2\pi}\sigma} \right) \left(e^{-\frac{y_2^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \right) \left(e^{-\frac{y_1^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\lambda} \right) \left(e^{-\frac{y_0^2}{2\lambda}} \right) \\
 \log_2 P(y_2, y_1, y_0 | \alpha) &= -\log_2(\sqrt{2\pi}\sigma) - \frac{(y_2 - \alpha y_1)^2}{2\sigma^2} - \log_2(\sqrt{2\pi}\sigma) - \frac{(y_1 - \alpha y_0)^2}{2\sigma^2} - \log_2(\sqrt{2\pi}\lambda) - \frac{y_0^2}{2\lambda} \\
 \downarrow \frac{\partial}{\partial \alpha} \\
 0 &= \frac{2y_1(y_2 - \alpha y_1)}{2\sigma^2} + \frac{2y_0(y_1 - \alpha y_0)}{2\sigma^2} \\
 0 &= y_1 y_2 - \alpha y_1^2 + y_0 y_1 - y_0^2 \\
 \alpha(y_1^2 + y_0^2) &= y_1 y_2 + y_0 y_1 \\
 \alpha &= \frac{y_1 y_2 + y_0 y_1}{y_1^2 + y_0^2}
 \end{aligned}$$

0.2.1 OT1

$$\begin{aligned}
 \text{OT1: } \log_2(P_{N+1}, P_N, P_{N-1} \dots | \alpha) \\
 &= \left[\prod_{i=1}^N \left(\frac{1}{\sqrt{2\pi}\sigma} \right) e^{-\frac{y_i^2}{2\sigma^2}} \right] \left(\frac{1}{\sqrt{2\pi}\sigma} \right) \left(e^{-\frac{y_0^2}{2\sigma^2}} \right) \\
 \log \downarrow \\
 &= \sum_{i=1}^N \left(-\log_2(\sqrt{2\pi}\sigma) - \frac{(y_{i+1} - \alpha y_i)^2}{2\sigma^2} \right) - \log_2(\sqrt{2\pi}\sigma) - \frac{y_0^2}{2\sigma^2} \\
 \frac{\partial}{\partial \alpha} \downarrow \\
 &= \sum_{i=1}^N \left(\frac{2y_i(y_{i+1} - \alpha y_i)}{2\sigma^2} \right) \\
 \alpha \sum_{i=1}^N y_i^2 &= \sum_{i=1}^N y_i y_{i+1} \Rightarrow \alpha = \frac{\sum_{i=1}^N y_i y_{i+1}}{\sum_{i=1}^N y_i^2} \#
 \end{aligned}$$

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import norm
```

```
[ ]: def find_decision_boundary(mean_a, sd_a, mean_b, sd_b, prior_a, prior_b):
    x_axis = np.arange(-10, 10, 0.0001)

    likelihood_a = norm.pdf(x_axis, loc=mean_a, scale=sd_a)
    likelihood_b = norm.pdf(x_axis, loc=mean_b, scale=sd_b)

    evidence = likelihood_a * prior_a + likelihood_b * prior_b

    plt.plot(x_axis, (likelihood_a / evidence) * prior_a, label="posterior a")
    plt.plot(x_axis, (likelihood_b / evidence) * prior_b, label="posterior b")
    plt.legend()
    plt.show()

    likelihood_ratio = likelihood_a / likelihood_b
    priors_ratio = prior_b / prior_a


    boundary_position = x_axis[likelihood_ratio >= priors_ratio][0]
    decision_boundary = likelihood_b[x_axis == boundary_position]


    plt.plot(x_axis, likelihood_a, label="likelihood a")
    plt.plot(x_axis, likelihood_b, label="likelihood b")
    plt.axvline(x=boundary_position)
    plt.legend()
    plt.show()
```

```
return (boundary_position, decision_boundary[0])
```

0.3 T2

```
[ ]: print("Decision Boundary is at", find_decision_boundary(4, np.sqrt(2), 0, np.  
↪sqrt(2), 0.5, 0.5))
```







Decision Boundary is at (2.000099999972033, 0.10376649693006622)

0.4 T3

```
[ ]: print("Decision Boundary is at", find_decision_boundary(4, np.sqrt(2), 0, np.
      ↪sqrt(2), 0.75, 0.25))
```




Decision Boundary is at (1.4506999999733132, 0.1666857883072887)


0.5 OT2

$$\begin{aligned}
 \text{OT2: } P(x|w_1)P(w_1) &= P(x|w_2)P(w_2) \\
 \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu_1)^2}{2\sigma^2}} &= \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu_2)^2}{2\sigma^2}} \\
 e^{\frac{(x-\mu_1)^2 - (x-\mu_2)^2}{2\sigma^2}} &= 1 \\
 \frac{(x-\mu_1)^2 - (x-\mu_2)^2}{2\sigma^2} &= 0 \\
 (\mu_1 - \mu_2)(2x - \mu_1 - \mu_2) &= 0 \\
 x &= \frac{\mu_1 + \mu_2}{2} \quad \#
 \end{aligned}$$

0.6 OT3

```
[ ]: print("Decision Boundary:", find_decision_boundary(4, np.sqrt(2), 0, 2, 0.5, 0.
↪5))
```





Decision Boundary: (2.1031999999717925, 0.11474798458217128)

1 Employee Attrition Prediction

1.0.1 read CSV

```
[ ]: df = pd.read_csv('hr-employee-attrition-with-null.csv')
```

1.0.2 Dataset statistic

```
[ ]: df.describe()
```

```
[ ]:
      Unnamed: 0      Age  DailyRate  DistanceFromHome  Education \
count  1470.000000  1176.000000  1176.000000      1176.00000  1176.000000
mean    734.500000   37.134354   798.875850         9.37500    2.920918
std     424.496761    9.190317   406.957684         8.23049    1.028796
min         0.000000   18.000000   102.000000         1.00000    1.000000
25%     367.250000   30.000000   457.750000         2.00000    2.000000
50%     734.500000   36.000000   798.500000         7.00000    3.000000
75%    1101.750000   43.000000  1168.250000        15.00000    4.000000
```

max	1469.000000	60.000000	1499.000000	29.00000	5.000000
-----	-------------	-----------	-------------	----------	----------

	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	HourlyRate \
count	1176.0	1176.000000	1176.000000	1176.000000
mean	1.0	1031.399660	2.733844	65.821429
std	0.0	601.188955	1.092992	20.317323
min	1.0	1.000000	1.000000	30.000000
25%	1.0	494.750000	2.000000	48.000000
50%	1.0	1027.500000	3.000000	66.000000
75%	1.0	1562.250000	4.000000	84.000000
max	1.0	2068.000000	4.000000	100.000000

	JobInvolvement	...	RelationshipSatisfaction	StandardHours \
count	1176.000000	...	1176.000000	1176.0
mean	2.728741	...	2.694728	80.0
std	0.705280	...	1.093660	0.0
min	1.000000	...	1.000000	80.0
25%	2.000000	...	2.000000	80.0
50%	3.000000	...	3.000000	80.0
75%	3.000000	...	4.000000	80.0
max	4.000000	...	4.000000	80.0

	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear \
count	1176.000000	1176.000000	1176.000000
mean	0.752551	11.295068	2.787415
std	0.822550	7.783376	1.290507
min	0.000000	0.000000	0.000000
25%	0.000000	6.000000	2.000000
50%	1.000000	10.000000	3.000000
75%	1.000000	15.000000	3.000000
max	3.000000	40.000000	6.000000

	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole \
count	1176.000000	1176.000000	1176.000000
mean	2.770408	7.067177	4.290816
std	0.705004	6.127836	3.630901
min	1.000000	0.000000	0.000000
25%	2.000000	3.000000	2.000000
50%	3.000000	5.000000	3.000000
75%	3.000000	10.000000	7.000000
max	4.000000	37.000000	18.000000

	YearsSinceLastPromotion	YearsWithCurrManager
count	1176.000000	1176.000000
mean	2.159014	4.096939
std	3.163524	3.537393
min	0.000000	0.000000

25%	0.000000	2.000000
50%	1.000000	3.000000
75%	2.250000	7.000000
max	15.000000	17.000000

[8 rows x 27 columns]

```
[ ]: df.head()
```

```
[ ]:      Unnamed: 0   Age Attrition      BusinessTravel  DailyRate \
0           0  41.0      Yes      Travel_Rarely      NaN
1           1   NaN      No              NaN      279.0
2           2  37.0      Yes              NaN     1373.0
3           3   NaN      No  Travel_Frequently     1392.0
4           4  27.0      No      Travel_Rarely     591.0

      Department  DistanceFromHome  Education  EducationField \
0           NaN              1.0      NaN  Life Sciences
1  Research & Development          NaN      NaN  Life Sciences
2           NaN              2.0      2.0      NaN
3  Research & Development          3.0      4.0  Life Sciences
4  Research & Development          2.0      1.0      Medical

      EmployeeCount  ...  RelationshipSatisfaction  StandardHours \
0           1.0  ...              1.0              80.0
1           1.0  ...              4.0              NaN
2           1.0  ...              NaN              80.0
3           NaN  ...              3.0              NaN
4           1.0  ...              4.0              80.0

      StockOptionLevel  TotalWorkingYears  TrainingTimesLastYear  WorkLifeBalance \
0           0.0              8.0              0.0              NaN
1           1.0             10.0              NaN              3.0
2           0.0              7.0              3.0              NaN
3           NaN              8.0              3.0              NaN
4           1.0              6.0              NaN              3.0

      YearsAtCompany  YearsInCurrentRole  YearsSinceLastPromotion \
0           6.0              NaN              0.0
1          10.0              NaN              NaN
2           NaN              0.0              NaN
3           8.0              NaN              3.0
4           2.0              2.0              2.0

      YearsWithCurrManager
0           NaN
1           7.0
```

```

2          0.0
3          0.0
4          NaN

```

[5 rows x 36 columns]

1.0.3 Feature transformation

```

[ ]: df.loc[df["Attrition"] == "no", "Attrition"] = 0.0
df.loc[df["Attrition"] == "yes", "Attrition"] = 1.0
string_categorical_col = ['Department', 'Attrition', 'BusinessTravel',
    ↪ 'EducationField', 'Gender', 'JobRole',
    ↪ 'MaritalStatus', 'Over18', 'OverTime']

# ENCODE STRING COLUMNS TO CATEGORICAL COLUMNS
for col in string_categorical_col:
    df[col] = pd.Categorical(df[col]).codes
    # INSERT CODE HERE

# HANDLE NULL NUMBERS
# df[np.isnan(df)] = -1
# INSERT CODE HERE

df = df.loc[:, ~df.columns.isin(['EmployeeNumber', 'Unnamed: 0',
    ↪ 'EmployeeCount', 'StandardHours', 'Over18'])]

```

1.0.4 Splitting data into train and test

```

[ ]: from sklearn.model_selection import train_test_split

[ ]: df_train, df_test = train_test_split(df, test_size=0.1, train_size=0.9,
    ↪ stratify=df["Attrition"])

```

1.0.5 Display histogram of each feature

```

[ ]: def display_histogram(df, col_name, cls, n_bin = 40):
    train_col_no_nan = df[col_name][~(np.isnan(df[col_name]))]
    hist, bin_edge = np.histogram(train_col_no_nan, n_bin)

    hist = np.array(hist, dtype=float) / np.sum(hist)

    plt.fill_between(bin_edge.repeat(2)[1:-1], hist.repeat(2),
    ↪ facecolor="steelblue")
    plt.title(cls)
    plt.show()


    print("Zero bins:", np.sum(hist == 0))

```


```
# INSERT CODE HERE
```

1.0.6 T4. Observe the histogram for Age, MonthlyIncome and DistanceFromHome. How many bins have zero counts? Do you think this is a good discretization? Why?

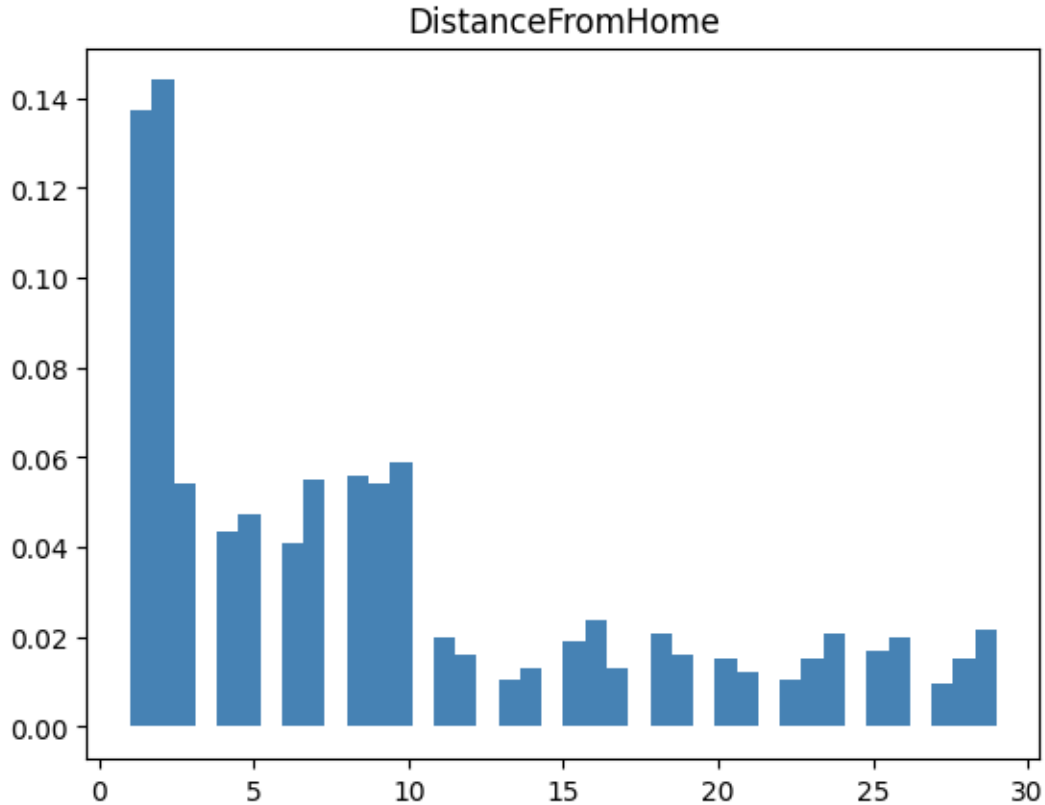
```
[ ]: display_histogram(df_train, "Age", "Age", 40)
      display_histogram(df_train, "MonthlyIncome", "MonthlyIncome", 40)
      display_histogram(df_train, "DistanceFromHome", "DistanceFromHome", 40)
```



Zero bins: 0



Zero bins: 1



Zero bins: 11

- Age has 0 zero bin
- MonthlyIncome has 0 zero bin
- DistanceFromHome has 11 zero bins

Feature `DistanceFromHome` does not get a good discretization with bin size 40 because it appears zero bins in the histogram even though `Age` and `MonthlyIncome` gets a smooth histogram

1.0.7 T5. Can we use a Gaussian to estimate this histogram? Why? What about a Gaussian Mixture Model (GMM)?

We are likely able to use Gaussian to estimate `Age` because the shape does not look like a normal gaussian distribution but for `MonthlyIncome`, it looks like two gaussian distribution in histogram which can be formulated as Gaussian Mixture Model but it's not likely to be used on `DistanceFromHome` because its shape look like logarithm graph instead

1.0.8 T6. Now plot the histogram according to the method described above (with 10, 40, and 100 bins) and show 3 plots each for Age, MonthlyIncome, and DistanceFromHome. Which bin size is most sensible for each features? Why?

```
[ ]: def calculate_bins_and_edges(arr, n_bin = 10):
    min_bin = -1e9
    max_bin = 1e9

    max_train = np.max(arr)
    min_train = np.min(arr)

    edges = np.linspace(start=min_train, stop=max_train, num=n_bin - 1)
    edges = np.insert(edges, 0, min_bin)
    edges = np.append(edges, max_bin)

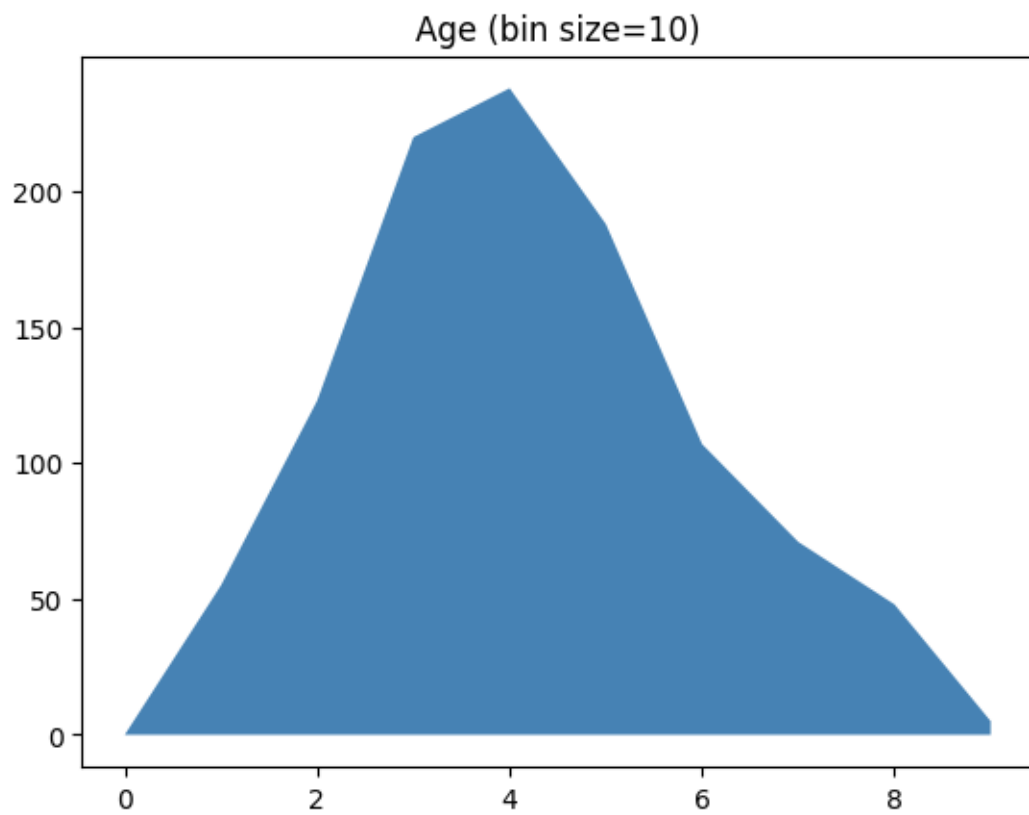
    digitized_df = np.digitize(arr, edges)
    bins = np.bincount(digitized_df)[1:]

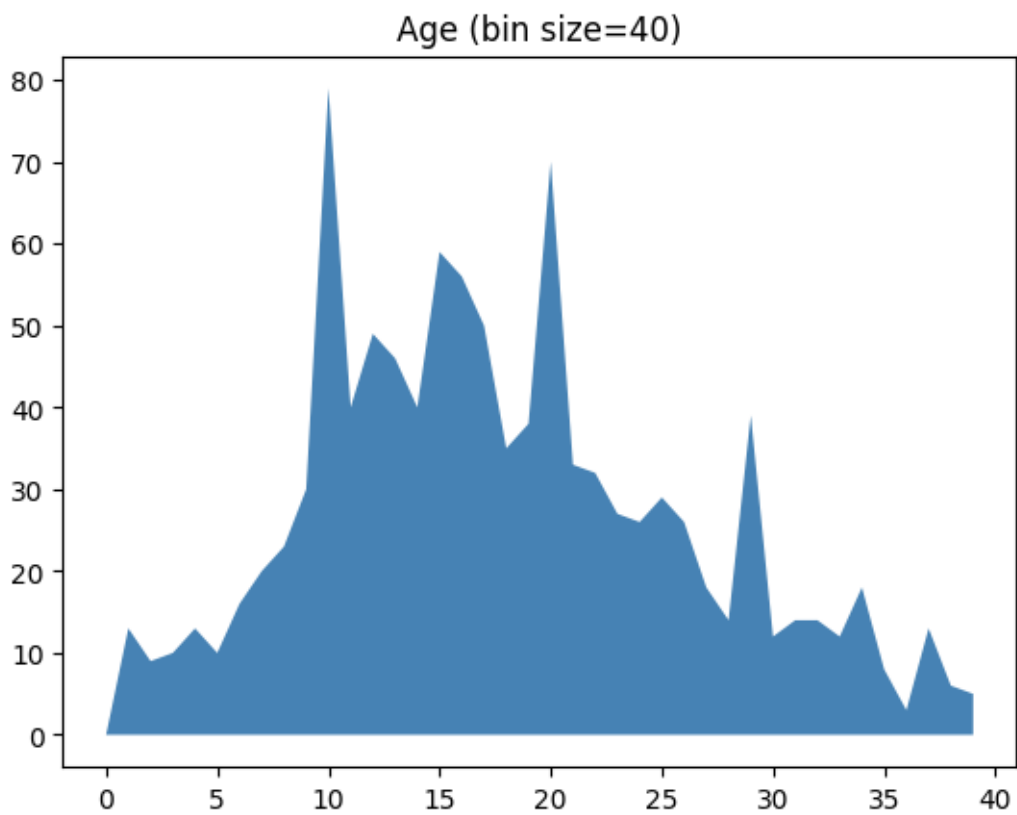
    return bins, edges


def plot_histogram_v2(df, col_name, n_bin = 10):
    train_col_no_nan = df[col_name][~(np.isnan(df[col_name]))]
    bins, _ = calculate_bins_and_edges(train_col_no_nan, n_bin)

    plt.fill_between(np.arange(0, n_bin), bins, facecolor="steelblue")
    plt.title(f"{col_name} (bin size={n_bin})")
    plt.show()

[ ]: plot_histogram_v2(df_train, "Age", 10)
     plot_histogram_v2(df_train, "Age", 40)
     plot_histogram_v2(df_train, "Age", 100)
```










For age feature, we will see that when bin size is 100 there start to have zero bins appearing in histogram but it look like a normal distribution curve also

```
[ ]: plot_histogram_v2(df_train, "MonthlyIncome", 10)
      plot_histogram_v2(df_train, "MonthlyIncome", 40)
      plot_histogram_v2(df_train, "MonthlyIncome", 100)
```










For MonthlyIncome feature, we will see that the sensible point start at bin size 100 which start to see the zero bins that should not appear in probabilistic estimation

```
[ ]: plot_histogram_v2(df_train, "DistanceFromHome", 10)
      plot_histogram_v2(df_train, "DistanceFromHome", 40)
      plot_histogram_v2(df_train, "DistanceFromHome", 100)
```







For DistanceFromHome feature, we will see that at least bin size 40 it seems to appear the zero bins in histogram

1.0.9 T7. For the rest of the features, which one should be discretized in order to be modeled by histograms? What are the criteria for choosing whether we should discretize a feature or not? Answer this and discretize those features into 10 bins each. In other words, figure out the bin edge for each feature, then use `digitize()` to convert the features to discrete values

```
[ ]: df
```

```
[ ]:
   Age  Attrition  BusinessTravel  DailyRate  Department  \
0   41.0         1                2         NaN          -1
1    NaN         0               -1        279.0           1
2   37.0         1               -1       1373.0          -1
3    NaN         0                1       1392.0           1
4   27.0         0                2        591.0           1
...   ...       ...              ...       ...       ...
1465  36.0         0                1        884.0           1
1466  39.0         0                2        613.0          -1
1467  27.0         0               -1        155.0           1
```

1468	49.0	0	1	1023.0	2
1469	34.0	0	2	628.0	1

	DistanceFromHome	Education	EducationField	EnvironmentSatisfaction	\
0	1.0	NaN	1	2.0	
1	NaN	NaN	1	3.0	
2	2.0	2.0	-1	NaN	
3	3.0	4.0	1	NaN	
4	2.0	1.0	3	1.0	
...	
1465	NaN	NaN	3	3.0	
1466	6.0	NaN	3	4.0	
1467	4.0	3.0	1	2.0	
1468	2.0	3.0	-1	4.0	
1469	NaN	NaN	-1	2.0	

	Gender	...	PerformanceRating	RelationshipSatisfaction	\
0	0	...	NaN	1.0	
1	1	...	NaN	4.0	
2	1	...	3.0	NaN	
3	0	...	3.0	3.0	
4	1	...	3.0	4.0	
...	
1465	1	...	3.0	3.0	
1466	-1	...	3.0	NaN	
1467	1	...	NaN	2.0	
1468	1	...	3.0	4.0	
1469	1	...	3.0	NaN	

	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear	\
0	0.0	8.0	0.0	
1	1.0	10.0	NaN	
2	0.0	7.0	3.0	
3	NaN	8.0	3.0	
4	1.0	6.0	NaN	
...	
1465	1.0	17.0	3.0	
1466	NaN	9.0	5.0	
1467	1.0	6.0	0.0	
1468	0.0	17.0	NaN	
1469	0.0	6.0	3.0	

	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole	\
0	NaN	6.0	NaN	
1	3.0	10.0	NaN	
2	NaN	NaN	0.0	
3	NaN	8.0	NaN	

4	3.0	2.0	2.0
...
1465	3.0	5.0	2.0
1466	3.0	7.0	7.0
1467	3.0	6.0	NaN
1468	2.0	9.0	6.0
1469	4.0	4.0	NaN

	YearsSinceLastPromotion	YearsWithCurrManager
0	0.0	NaN
1	NaN	7.0
2	NaN	0.0
3	3.0	0.0
4	2.0	NaN
...
1465	0.0	3.0
1466	1.0	7.0
1467	0.0	3.0
1468	0.0	8.0
1469	1.0	2.0


[1470 rows x 31 columns]

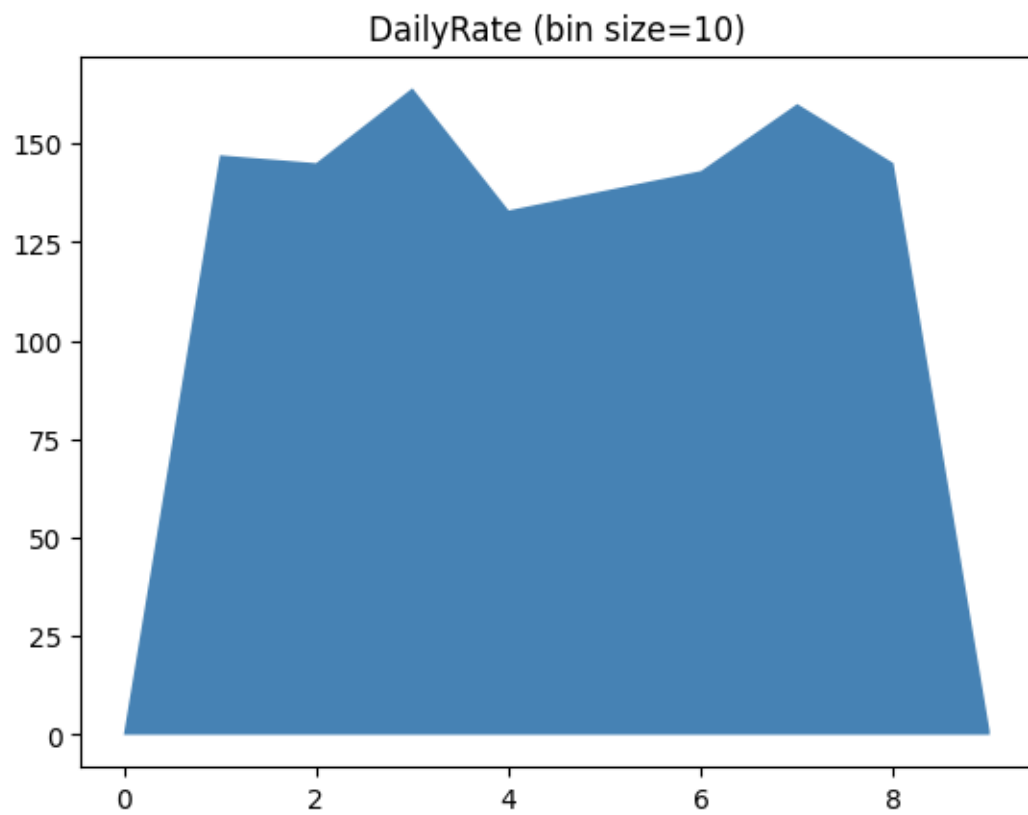
I will use the continuity of histogram as a criteria for considering which features should be discretized


- Age
- DailyRate
- DistanceFromHome
- HourlyRate
- JobRole
- MonthlyIncome
- MonthlyRate
- NumCompaniesWorked
- PercentSalaryHike
- TotalWorkingYears
- TrainingTimesLastYear
- YearsAtCompany
- YearsInCurrentRole
- YearsSinceLastPromotion
- YearsWithCurrManager

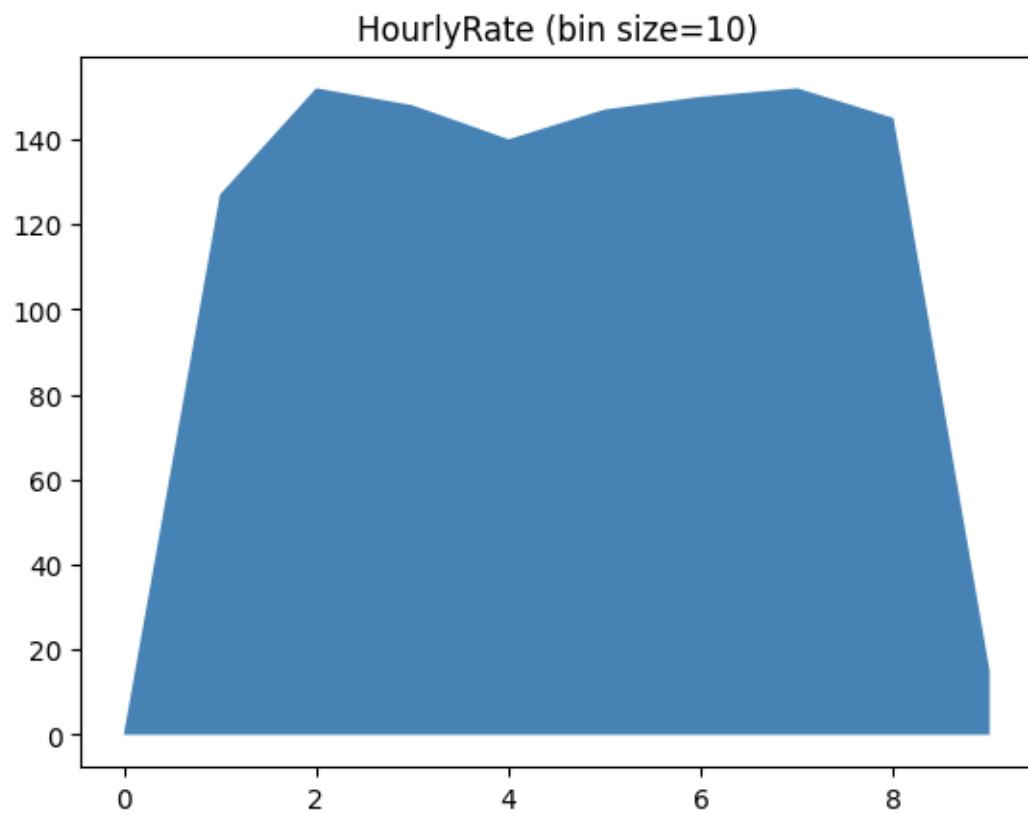
```
[ ]: column_name = ["Age", "DailyRate", "DistanceFromHome", "HourlyRate", "JobRole",
    ↪ "MonthlyIncome", "MonthlyRate", "NumCompaniesWorked", "PercentSalaryHike",
    ↪ "TotalWorkingYears", "TrainingTimesLastYear", "YearsAtCompany",
    ↪ "YearsInCurrentRole", "YearsSinceLastPromotion", "YearsWithCurrManager"]

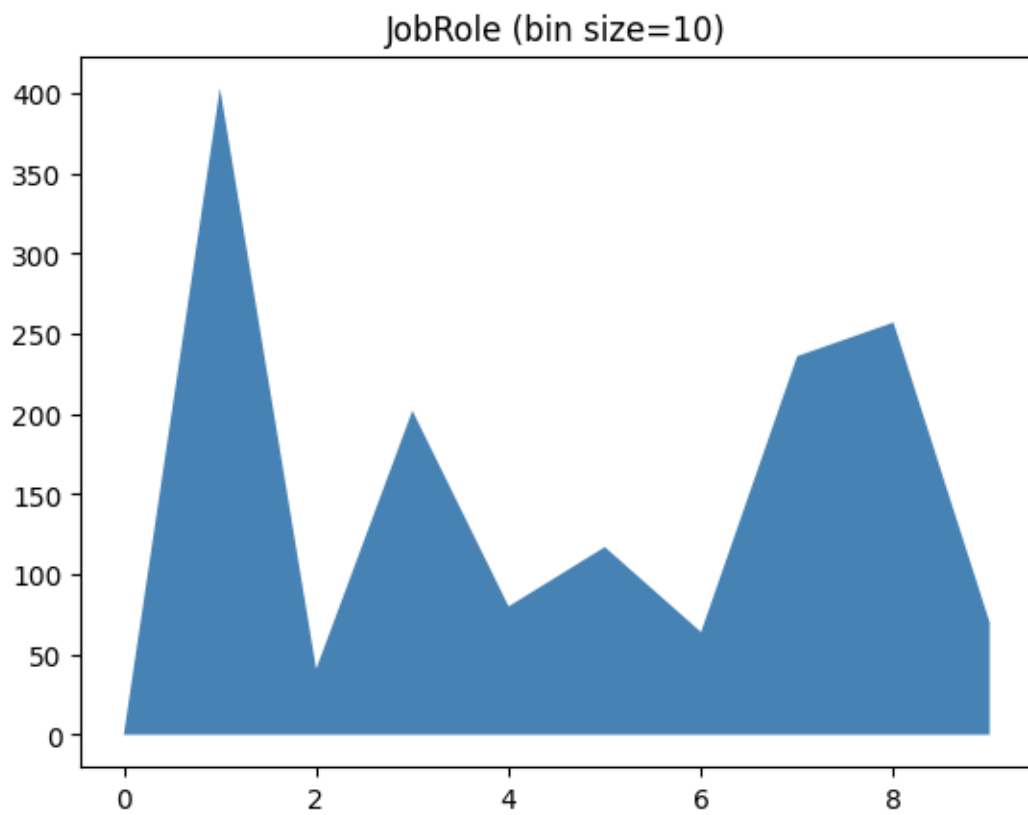
for col in column_name:
    plot_histogram_v2(df, col, 10)
```

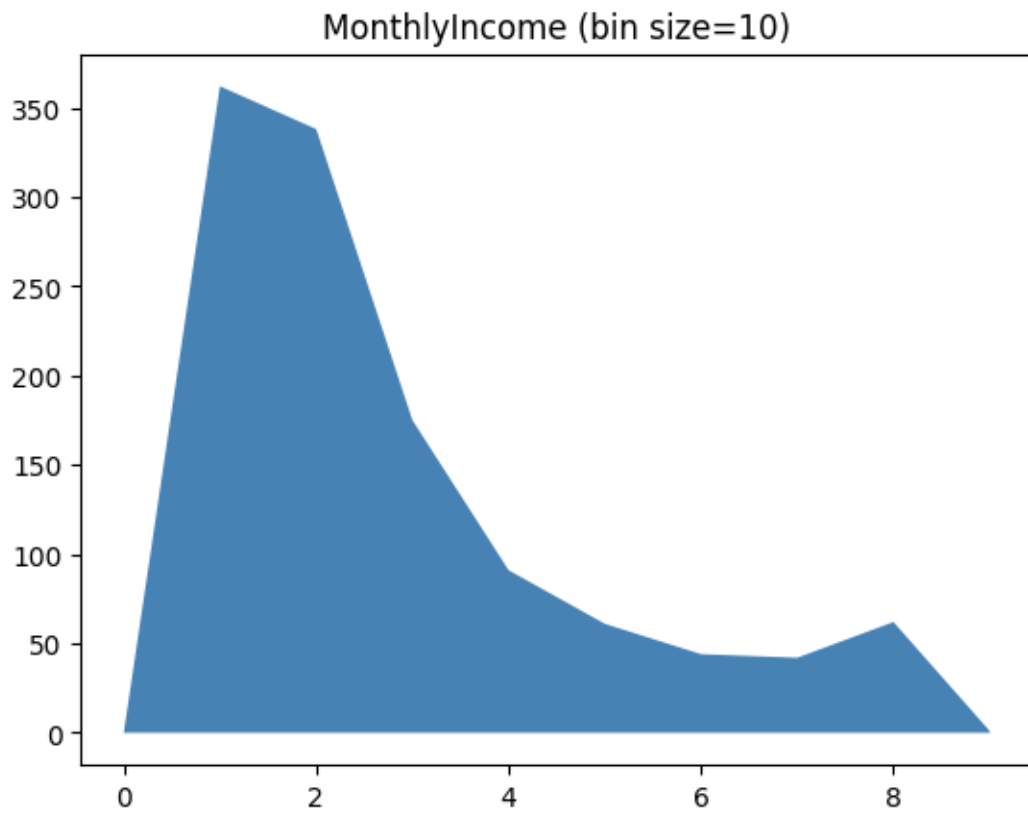


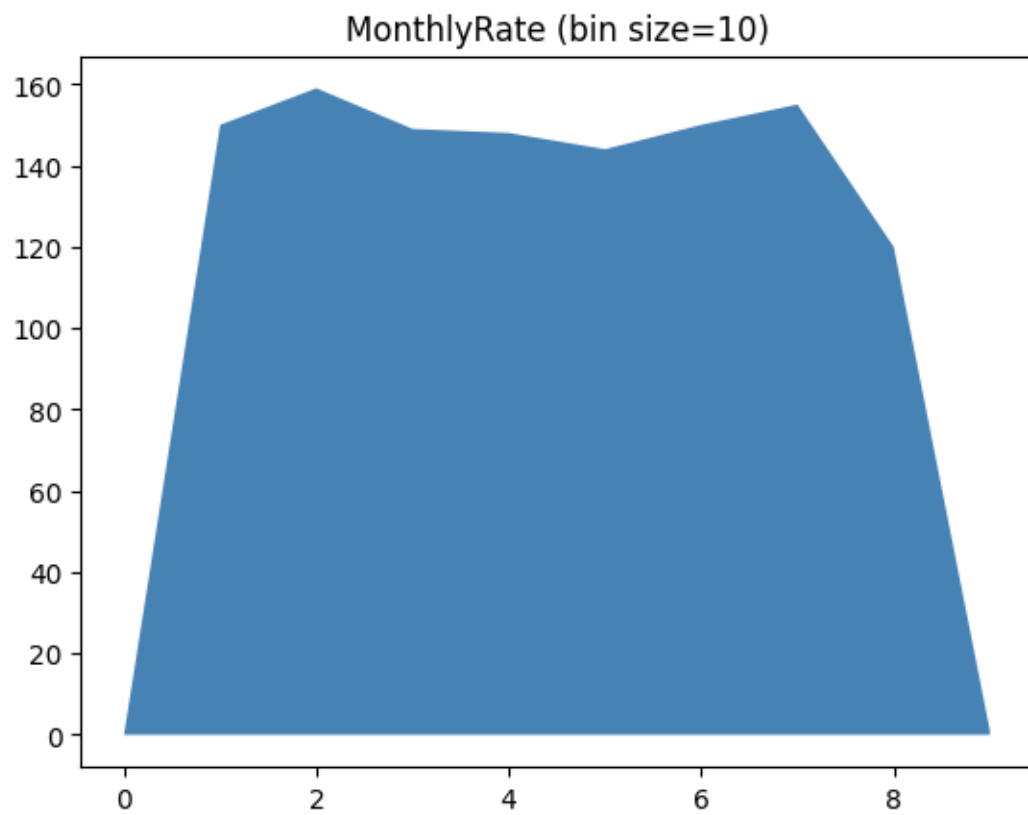


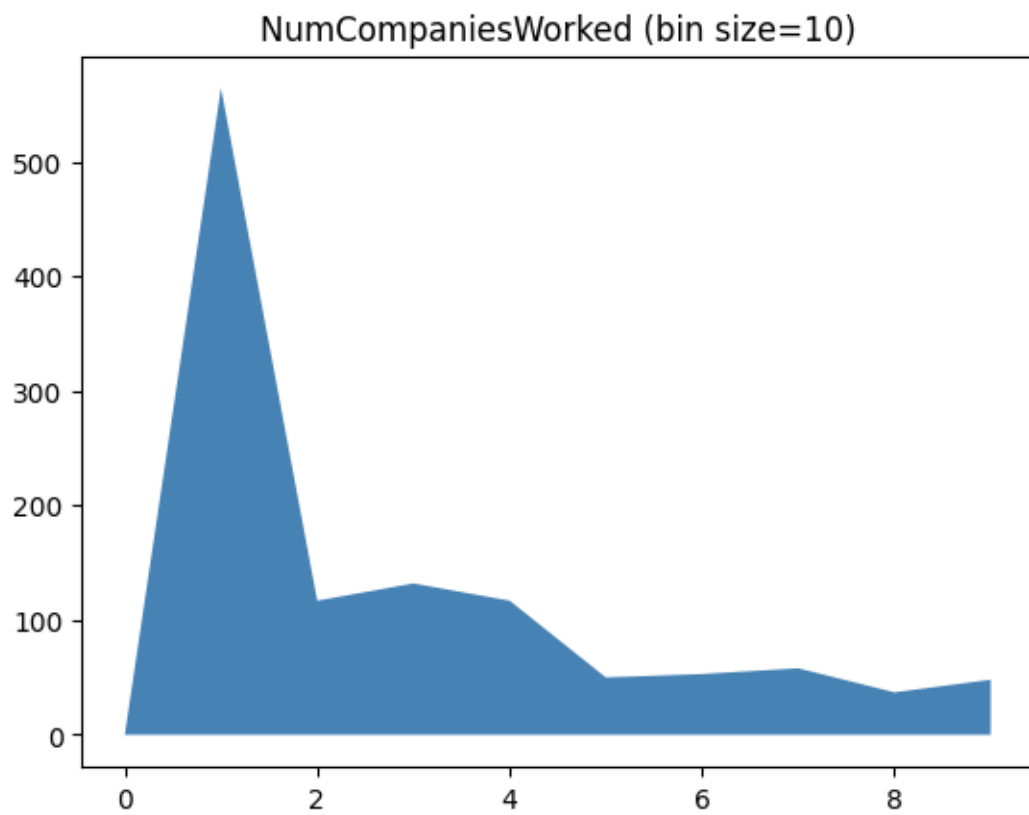


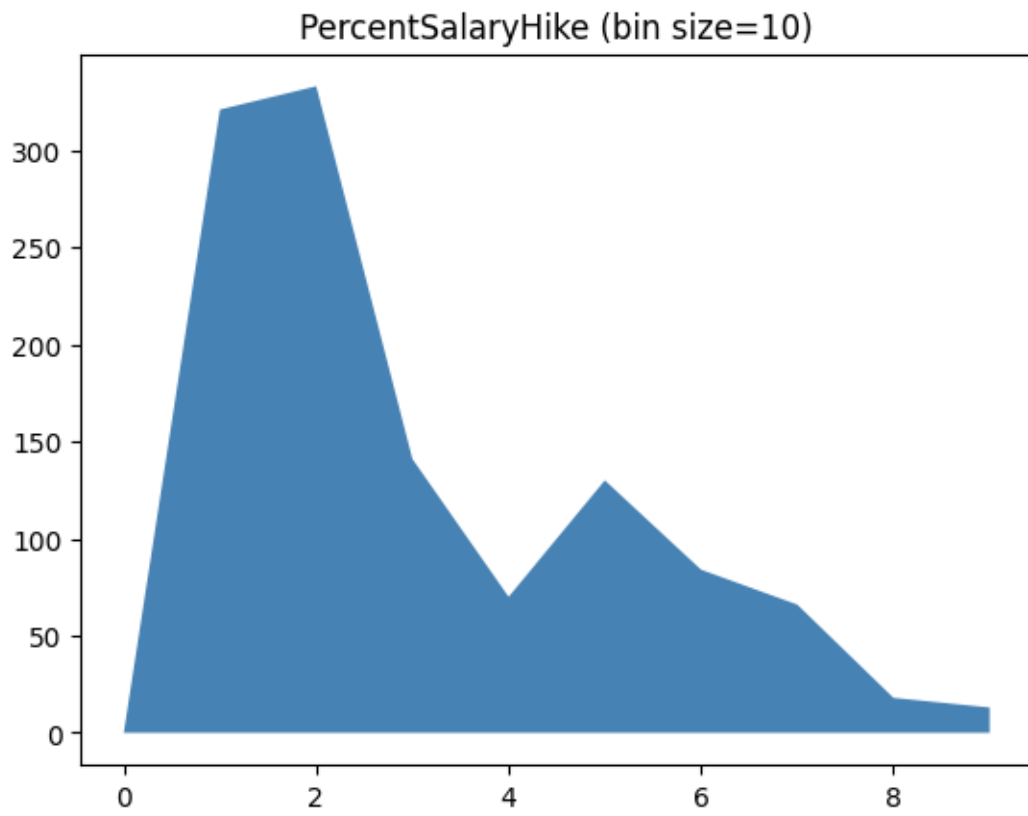


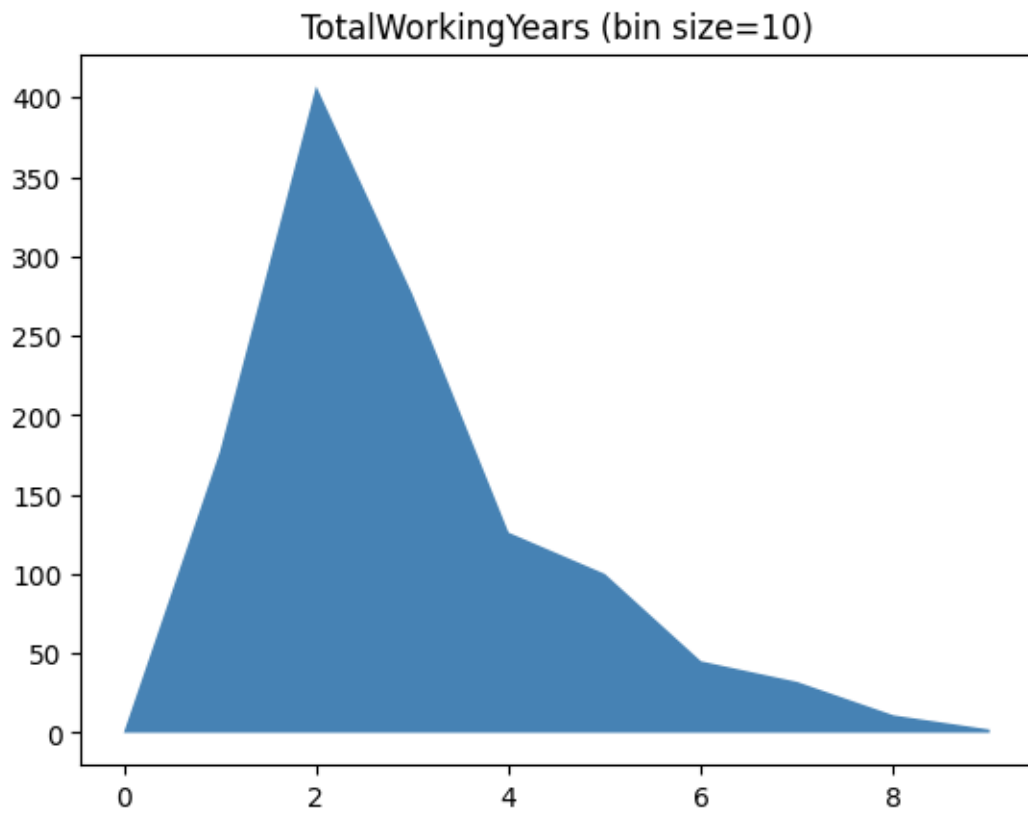


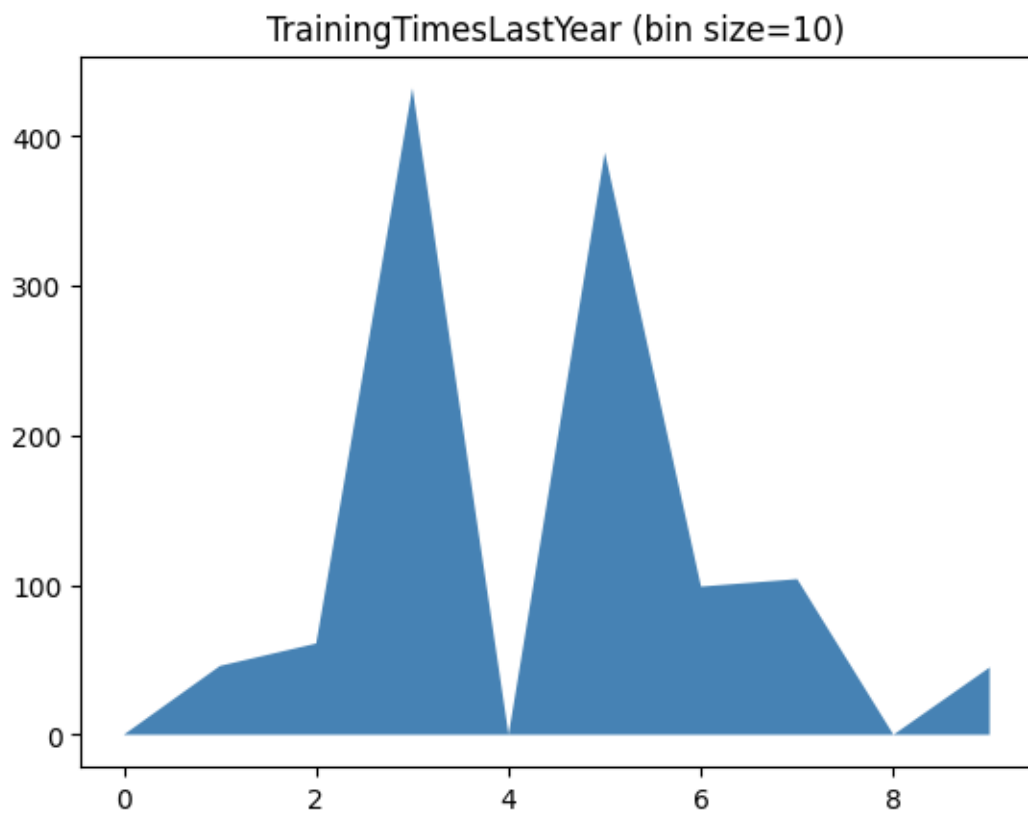


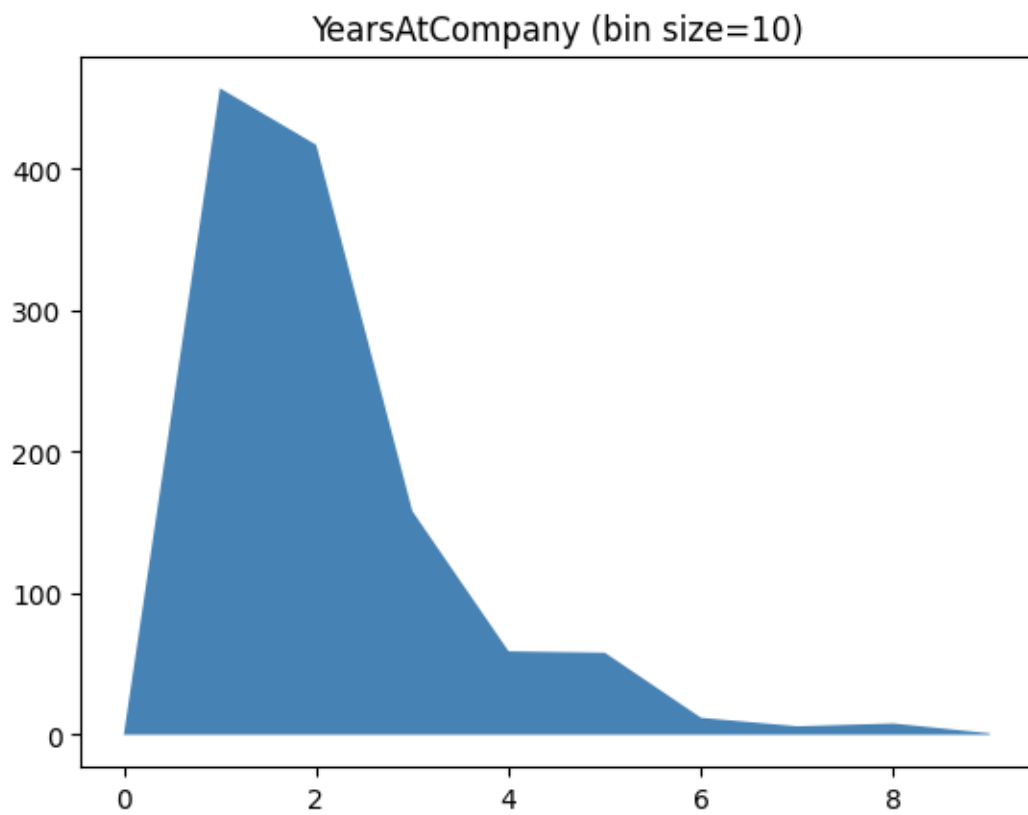


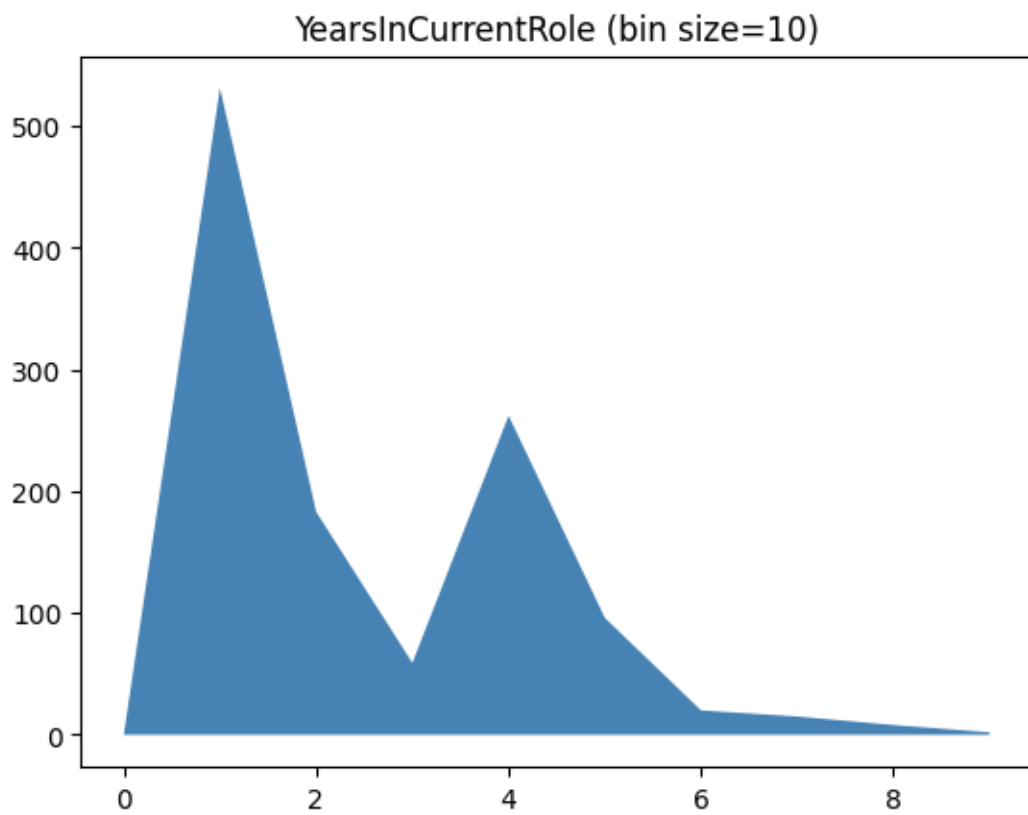


















1.0.10 T8. What kind of distribution should we use to model histograms? (Answer a distribution name) What is the MLE for the likelihood distribution? (Describe how to do the MLE). Plot the likelihood distributions of MonthlyIncome, JobRole, HourlyRate, and MaritalStatus for different Attrition values.

The distribution which we will use for describing histograms is categorical distribution which describe each bins as parameter and x_i as frequency in each bins

$$P(X = x) = p_1^{x_1} p_2^{x_2} \dots p_k^{x_k}$$

We will find p_i which maximize $P(X = x)$ by using lagrange multiplier

$$\begin{aligned}
 \text{TS: } L(\theta) &= \prod_{i=1}^N \theta_i^{x_i} \\
 \log L(\theta) &= \sum_{i=1}^N x_i \log \theta_i \\
 f(\theta_1, \theta_2, \dots, \theta_n) &= \log L(\theta) \rightarrow \nabla f(\theta_1, \theta_2, \theta_3, \dots, \theta_n) = \left(\frac{x_1}{\theta_1}, \frac{x_2}{\theta_2}, \dots, \frac{x_n}{\theta_n} \right) \\
 g(\theta_1, \theta_2, \dots, \theta_n) &= \sum_{i=1}^N \theta_i \rightarrow \nabla g(\theta_1, \theta_2, \theta_3, \dots, \theta_n) = (1, 1, 1, \dots, 1) \\
 \text{Solve for } \nabla f(\theta_1, \theta_2, \dots, \theta_n) &= \lambda \nabla g(\theta_1, \theta_2, \dots, \theta_n) \\
 \frac{x_1}{\theta_1} &= \lambda, \frac{x_2}{\theta_2} = \lambda \dots \frac{x_n}{\theta_n} = \lambda \\
 \therefore x_1 &= \lambda \theta_1, x_2 = \lambda \theta_2, \dots, x_n = \lambda \theta_n \\
 g(\theta_1, \theta_2, \dots, \theta_n) &= \sum_{i=1}^N \theta_i = \frac{x_1}{\lambda} + \frac{x_2}{\lambda} + \frac{x_3}{\lambda} + \dots + \frac{x_n}{\lambda} = 1 \\
 \therefore \lambda &= x_1 + x_2 + x_3 + \dots + x_n \\
 \therefore \theta_1 &= \frac{x_1}{\sum x_i}, \theta_2 = \frac{x_2}{\sum x_i}, \dots, \theta_n = \frac{x_n}{\sum x_i}
 \end{aligned}$$


```
[ ]: def plot_likelihood_attrition(df, y, attrition, col_name, n_bin=10):
    features = df[col_name][~(np.isnan(df[col_name]))]
    features = features[y == attrition]


    bins, edges = calculate_bins_and_edges(features, n_bin)
    bins = bins / np.sum(bins)

    plt.fill_between(np.arange(0, n_bin), bins, facecolor="steelblue")
    plt.title(f"{col_name} (attrition={attrition})")
    plt.show()
```


```
[ ]: y_train = df_train["Attrition"]
```


```
[ ]: plot_likelihood_attrition(df_train, y_train, 0, "MonthlyIncome")
plot_likelihood_attrition(df_train, y_train, 1, "MonthlyIncome")
```







```
[ ]: plot_likelihood_attrition(df_train, y_train, 0, "JobRole")  
      plot_likelihood_attrition(df_train, y_train, 1, "JobRole")
```







```
[ ]: plot_likelihood_attrition(df_train, y_train, 0, "HourlyRate")  
      plot_likelihood_attrition(df_train, y_train, 1, "HourlyRate")
```





```
[ ]: plot_likelihood_attrition(df_train, y_train, 0, "MaritalStatus")  
plot_likelihood_attrition(df_train, y_train, 1, "MaritalStatus")
```





1.0.11 T9. What is the prior distribution of the two classes?

$$P(\theta) = (0.16)^\theta * (0.84)^{1-\theta}; \theta = 0, 1$$

```
[ ]: n_stay = np.sum(y_train == 1)
      n_leave = np.sum(y_train == 0)

      prior_stay = n_stay / (n_stay + n_leave)
      prior_leave = n_leave / (n_stay + n_leave)

      print(f"Prior class stay: {prior_stay}")
      print(f"Prior class leave: {prior_leave}")
```

Prior class stay: 0.16099773242630386

Prior class leave: 0.8390022675736961

1.0.12 T10. If we use the current Naive Bayes with our current Maximum Likelihood Estimates, we will find that some $P(x_i | attrition)$ will be zero and will result in the entire product term to be zero. Propose a method to fix this problem.

I think one of the method to prevent zero probability is to distribute a small value from non-zero term to add to its, this way will not affect to overall probability much and still preserve zero

probability

1.0.13 T11. Implement your Naive Bayes classifier. Use the learned distributions to classify the test set. Don't forget to allow your classifier to handle missing values in the test set. Report the overall Accuracy. Then, report the Precision, Recall, and F score for detecting attrition. See Lecture 1 for the definitions of each metric.

```
[ ]: import random as rnd
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from scipy import stats

class SimpleBayesClassifier:

    def __init__(self, n_pos, n_neg):

        """
        Initializes the SimpleBayesClassifier with prior probabilities.

        Parameters:
        n_pos (int): The number of positive samples.
        n_neg (int): The number of negative samples.

        Returns:
        None: This method does not return anything as it is a constructor.
        """

        self.n_pos = n_pos
        self.n_neg = n_neg
        self.prior_pos = (n_pos) / (n_pos + n_neg)
        self.prior_neg = (n_neg) / (n_pos + n_neg)

    def fit_params(self, x, y, n_bins = 10):

        """
        Computes histogram-based parameters for each feature in the dataset.

        Parameters:
        x (np.ndarray): The feature matrix, where rows are samples and columns
        are features.
        y (np.ndarray): The target array, where each element corresponds to the
        label of a sample.
        n_bins (int): Number of bins to use for histogram calculation.

        Returns:
```

(stay_params, leave_params): A tuple containing two lists of tuples, one for 'stay' parameters and one for 'leave' parameters. Each tuple in the list contains the bins and edges of the histogram for a feature.

"""

```

self.stay_params = [(None, None) for _ in range(x.shape[1])]
self.leave_params = [(None, None) for _ in range(x.shape[1])]

stay_df = x[y == 1]
leave_df = x[y == 0]

column_width = x.shape[1]

for i in range(column_width):
    feature_stay_df = stay_df[:, i]
    feature_stay_df = feature_stay_df[~(np.isnan(feature_stay_df))]

    bins, edges = calculate_bins_and_edges(feature_stay_df, n_bins)
    bins = bins / np.sum(bins)

    x = np.sum(bins == 0)

    bins[bins == 0] = 1e-6
    bins[bins.argmax()] -= 1e-6 * x

    self.stay_params[i] = (bins, edges)

for i in range(column_width):
    feature_leave_df = leave_df[:, i]
    feature_leave_df = feature_leave_df[~(np.isnan(feature_leave_df))]

    bins, edges = calculate_bins_and_edges(feature_leave_df, n_bins)
    bins = bins / np.sum(bins)

    x = np.sum(bins == 0)

    bins[bins == 0] = 1e-6
    bins[bins.argmax()] -= 1e-6 * x

    self.leave_params[i] = (bins, edges)

return self.stay_params, self.leave_params

def predict(self, x, thresh = 0):
    """

```

*Predicts the class labels for the given samples using the
↳non-parametric model.*

Parameters:

x (np.ndarray): The feature matrix for which predictions are to be made.

*thresh (float): The threshold for log probability to decide between
↳classes.*

Returns:

*result (list): A list of predicted class labels (0 or 1) for each
↳sample in the feature matrix.*

"""

```
y_pred = []
classifier_arr = []

for data in x:
    classifier = np.log(self.prior_neg) - np.log(self.prior_pos)
    for (idx, feature) in enumerate(data):
        if np.isnan(feature):
            continue

        bins_pos, edges_pos = self.stay_params[idx]

        digitalized_feature_pos = np.digitize(feature, edges_pos) - 1
        log_prob_pos = np.log(bins_pos[digitalized_feature_pos])

        bins_neg, edges_neg = self.leave_params[idx]

        digitalized_feature_neg = np.digitize(feature, edges_neg) - 1
        log_prob_neg = np.log(bins_neg[digitalized_feature_neg])

        classifier += log_prob_neg - log_prob_pos

    classifier_arr.append(classifier)
    y_pred.append(1 if classifier <= thresh else 0)

return np.array(y_pred), np.array(classifier_arr)
```

```
def fit_gaussian_params(self, x: np.ndarray, y: np.ndarray):
```

"""

Computes mean and standard deviation for each feature in the dataset.

Parameters:

*x (np.ndarray): The feature matrix, where rows are samples and columns
↳are features.*

y (np.ndarray): The target array, where each element corresponds to the label of a sample.

Returns:

(gaussian_stay_params, gaussian_leave_params): A tuple containing two lists of tuples, one for 'stay' parameters and one for 'leave' parameters. Each tuple in the list contains the mean and standard deviation for a feature.

```
"""  
  
self.gaussian_stay_params = [(0, 0) for _ in range(x.shape[1])]  
self.gaussian_leave_params = [(0, 0) for _ in range(x.shape[1])]  
  
stay_df = x[y == 1]  
leave_df = x[y == 0]  
  
column_width = x.shape[1]  
  
for i in range(column_width):  
    feature_stay_df = stay_df[:, i]  
    feature_stay_df = feature_stay_df[~(np.isnan(feature_stay_df))]  
  
    self.gaussian_stay_params[i] = (np.mean(feature_stay_df), np.  
std(feature_stay_df))  
  
    feature_leave_df = leave_df[:, i]  
    feature_leave_df = feature_leave_df[~(np.isnan(feature_leave_df))]  
  
    self.gaussian_leave_params[i] = (np.mean(feature_leave_df), np.  
std(feature_leave_df))  
  
return self.gaussian_stay_params, self.gaussian_leave_params  
  
def gaussian_predict(self, x:np.ndarray, thresh = 0):  
    """  
    Predicts the class labels for the given samples using the parametric  
model.  
  
    Parameters:  
    x (np.ndarray): The feature matrix for which predictions are to be made.  
    thresh (float): The threshold for log probability to decide between  
classes.
```

Returns:

result (list): A list of predicted class labels (0 or 1) for each sample in the feature matrix.

```

"""

y_pred = []
classifier_arr = []

# INSERT CODE HERE
for data in x:
    classifier = np.log(self.prior_neg) - np.log(self.prior_pos)
    for (idx, feature) in enumerate(data):
        if np.isnan(feature):
            continue

        log_prob_pos = stats.norm.logpdf(feature, self.
gaussian_stay_params[idx][0], self.gaussian_stay_params[idx][1])
        log_prob_neg = stats.norm.logpdf(feature, self.
gaussian_leave_params[idx][0], self.gaussian_leave_params[idx][1])

        classifier += log_prob_neg - log_prob_pos

    classifier_arr.append(classifier)
    y_pred.append(1 if classifier <= thresh else 0)

return np.array(y_pred), np.array(classifier_arr)

```

```

[ ]: data_train = df_train.to_numpy()
data_test = df_test.to_numpy()

df

```

```

[ ]:

```

	Age	Attrition	BusinessTravel	DailyRate	Department \
0	41.0	1	2	NaN	-1
1	NaN	0	-1	279.0	1
2	37.0	1	-1	1373.0	-1
3	NaN	0	1	1392.0	1
4	27.0	0	2	591.0	1
...
1465	36.0	0	1	884.0	1
1466	39.0	0	2	613.0	-1
1467	27.0	0	-1	155.0	1
1468	49.0	0	1	1023.0	2
1469	34.0	0	2	628.0	1

	DistanceFromHome	Education	EducationField	EnvironmentSatisfaction \
0	1.0	NaN	1	2.0
1	NaN	NaN	1	3.0

2	2.0	2.0	-1	NaN
3	3.0	4.0	1	NaN
4	2.0	1.0	3	1.0
...
1465	NaN	NaN	3	3.0
1466	6.0	NaN	3	4.0
1467	4.0	3.0	1	2.0
1468	2.0	3.0	-1	4.0
1469	NaN	NaN	-1	2.0

	Gender	...	PerformanceRating	RelationshipSatisfaction	\
0	0	...	NaN	1.0	
1	1	...	NaN	4.0	
2	1	...	3.0	NaN	
3	0	...	3.0	3.0	
4	1	...	3.0	4.0	
...	
1465	1	...	3.0	3.0	
1466	-1	...	3.0	NaN	
1467	1	...	NaN	2.0	
1468	1	...	3.0	4.0	
1469	1	...	3.0	NaN	

	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear	\
0	0.0	8.0	0.0	
1	1.0	10.0	NaN	
2	0.0	7.0	3.0	
3	NaN	8.0	3.0	
4	1.0	6.0	NaN	
...	
1465	1.0	17.0	3.0	
1466	NaN	9.0	5.0	
1467	1.0	6.0	0.0	
1468	0.0	17.0	NaN	
1469	0.0	6.0	3.0	

	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole	\
0	NaN	6.0	NaN	
1	3.0	10.0	NaN	
2	NaN	NaN	0.0	
3	NaN	8.0	NaN	
4	3.0	2.0	2.0	
...	
1465	3.0	5.0	2.0	
1466	3.0	7.0	7.0	
1467	3.0	6.0	NaN	
1468	2.0	9.0	6.0	

1469	4.0	4.0	NaN
------	-----	-----	-----

	YearsSinceLastPromotion	YearsWithCurrManager
0	0.0	NaN
1	NaN	7.0
2	NaN	0.0
3	3.0	0.0
4	2.0	NaN
...
1465	0.0	3.0
1466	1.0	7.0
1467	0.0	3.0
1468	0.0	8.0
1469	1.0	2.0

[1470 rows x 31 columns]

```
[ ]: x_train = np.delete(data_train, 1, axis=1)
      y_train = data_train[:,1]

      x_test = np.delete(data_test, 1, axis=1)
      y_test = data_test[:,1]
```

```
[ ]: model = SimpleBayesClassifier(n_pos = np.sum(y_train == 1), n_neg = np.
      ↪sum(y_train == 0))
```

```
[ ]: def check_prior():
      """
      This function designed to test the implementation of the prior probability_
      ↪calculation in a Naive Bayes classifier.
      Specifically, it checks if the classifier correctly computes the prior_
      ↪probabilities for the
      negative and positive classes based on given input counts.
      """

      # prior_neg = 5/(5 + 5) = 0.5 and # prior_pos = 5/(5 + 5) = 0.5
      assert (SimpleBayesClassifier(5, 5).prior_pos, SimpleBayesClassifier(5, 5).
      ↪prior_neg) == (0.5, 0.5)

      assert (SimpleBayesClassifier(3, 5).prior_pos, SimpleBayesClassifier(3, 5).
      ↪prior_neg) == (0.375, 0.625)
      assert (SimpleBayesClassifier(0, 1).prior_pos, SimpleBayesClassifier(0, 1).
      ↪prior_neg) == (0.0, 1.0)
      assert (SimpleBayesClassifier(1, 0).prior_pos, SimpleBayesClassifier(1, 0).
      ↪prior_neg) == (1.0, 0.0)

      check_prior()
```



```
[ ]: model.fit_params(x_train, y_train)
```

```
[ ]: def check_fit_params():  
  
    """  
    This function is designed to test the fit_params method of a  
    ↪ SimpleBayesClassifier.  
    This method is presumably responsible for computing parameters for a Naive  
    ↪ Bayes classifier  
    based on the provided training data. The parameters in this context is bins  
    ↪ and edges from each histogram.  
    """  
  
    T = SimpleBayesClassifier(2, 2)  
    X_TRAIN_CASE_1 = np.array([  
        [0, 1, 2, 3],  
        [1, 2, 3, 4],  
        [2, 3, 4, 5],  
        [3, 4, 5, 6]  
    ])  
    Y_TRAIN_CASE_1 = np.array([0, 1, 0, 1])  
    STAY_PARAMS_1, LEAVE_PARAMS_1 = T.fit_params(X_TRAIN_CASE_1, Y_TRAIN_CASE_1)  
  
    print("STAY PARAMETERS")  
    for f_idx in range(len(STAY_PARAMS_1)):  
        print(f"Feature : {f_idx}")  
        print(f"BINS : {STAY_PARAMS_1[f_idx][0]}")  
        print(f"EDGES : {STAY_PARAMS_1[f_idx][1]}")  
    print("")  
    print("LEAVE PARAMETERS")  
    for f_idx in range(len(STAY_PARAMS_1)):  
        print(f"Feature : {f_idx}")  
        print(f"BINS : {LEAVE_PARAMS_1[f_idx][0]}")  
        print(f"EDGES : {LEAVE_PARAMS_1[f_idx][1]}")  
  
    check_fit_params()
```

```
[ ]: y_pred, _ = model.predict(x = x_test)
```

```
[ ]: def evaluate(y_true, y_pred, cls = 0):  
    y_pred_arr = np.array(y_pred)  
    y_true_arr = np.array(y_true)  
  
    result = y_true_arr == y_pred_arr  
  
    accuracy = np.sum(result) / len(y_pred)  
    recall = np.sum(result[y_pred_arr == cls]) / np.sum(y_true_arr == cls)
```

```

precision = np.sum(result[y_pred_arr == cls]) / np.sum(y_pred_arr == cls)
F1 = 2 * (recall * precision) / (recall + precision)

fpr = np.sum(np.logical_and(y_true == 0, y_pred == 1)) / np.sum(y_true_arr == 0)

return accuracy, precision, recall, F1, fpr

```

```

[ ]: accuracy, precision, recall, F1, fpr = evaluate(y_test, y_pred, cls=1)

print(f"accuracy: {accuracy}")
print(f"precision: {precision}")
print(f"recall: {recall}")
print(f"F1 score: {F1}")

```

```

accuracy: 0.8707482993197279
precision: 0.6923076923076923
recall: 0.375
F1 score: 0.48648648648648646

```

1.0.14 T12. Use the learned distributions to classify the test set. Report the results using the same metric as the previous question.

```

[ ]: model.fit_gaussian_params(x_train, y_train)

```

```

[ ]: def check_fit_gaussian_params():

    """
    This function is designed to test the fit_gaussian_params method of a
    SimpleBayesClassifier.
    This method is presumably responsible for computing parameters for a Naive
    Bayes classifier
    based on the provided training data. The parameters in this context is mean
    and STD.
    """

    T = SimpleBayesClassifier(2, 2)
    X_TRAIN_CASE_1 = np.array([
        [0, 1, 2, 3],
        [1, 2, 3, 4],
        [2, 3, 4, 5],
        [3, 4, 5, 6]
    ])
    Y_TRAIN_CASE_1 = np.array([0, 1, 0, 1])
    STAY_PARAMS_1, LEAVE_PARAMS_1 = T.fit_gaussian_params(X_TRAIN_CASE_1,
    Y_TRAIN_CASE_1)

```

```

print("STAY PARAMETERS")
for f_idx in range(len(STAY_PARAMS_1)):
    print(f"Feature : {f_idx}")
    print(f"Mean : {STAY_PARAMS_1[f_idx][0]}")
    print(f"STD. : {STAY_PARAMS_1[f_idx][1]}")
print("")
print("LEAVE PARAMETERS")
for f_idx in range(len(STAY_PARAMS_1)):
    print(f"Feature : {f_idx}")
    print(f"Mean : {LEAVE_PARAMS_1[f_idx][0]}")
    print(f"STD. : {LEAVE_PARAMS_1[f_idx][1]}")

check_fit_gaussian_params()

```

STAY PARAMETERS

```

Feature : 0
Mean : 2.0
STD. : 1.0
Feature : 1
Mean : 3.0
STD. : 1.0
Feature : 2
Mean : 4.0
STD. : 1.0
Feature : 3
Mean : 5.0
STD. : 1.0

```

LEAVE PARAMETERS

```

Feature : 0
Mean : 1.0
STD. : 1.0
Feature : 1
Mean : 2.0
STD. : 1.0
Feature : 2
Mean : 3.0
STD. : 1.0
Feature : 3
Mean : 4.0
STD. : 1.0

```

```
[ ]: y_pred, _ = model.gaussian_predict(x_test)
```

```
[ ]: accuracy, precision, recall, F1, fpr = evaluate(y_test, y_pred, cls=0)
```

```

print(f"accuracy: {accuracy}")
print(f"precision: {precision}")

```

```
print(f"recall: {recall}")
print(f"F1 score: {F1}")
```

```
accuracy: 0.8367346938775511
precision: 0.889763779527559
recall: 0.9186991869918699
F1 score: 0.904
```

```
[ ]: accuracy, precision, recall, F1, fpr = evaluate(y_test, y_pred, cls=1)
```

```
print(f"accuracy: {accuracy}")
print(f"precision: {precision}")
print(f"recall: {recall}")
print(f"F1 score: {F1}")
```

```
accuracy: 0.8367346938775511
precision: 0.5
recall: 0.4166666666666667
F1 score: 0.45454545454545453
```

1.0.15 T13 : The random choice baseline is the accuracy if you make a random guess for each test sample. Give random guess (50% leaving, and 50% staying) to the test samples. Report the overall Accuracy. Then, report the Precision, Recall, and F score for attrition prediction using the random choice baseline.

```
[ ]: test_size = y_test.shape[0]

y_random_result = np.random.choice([0, 1], test_size)

y_random_result
```

```
[ ]: array([1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0,
          0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0,
          0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1,
          1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1,
          0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 1,
          0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1,
          0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0])
```

```
[ ]: accuracy, precision, recall, F1, fpr = evaluate(y_test, y_random_result, cls=1)

print(f"accuracy: {accuracy}")
print(f"precision: {precision}")
print(f"recall: {recall}")
print(f"F1 score: {F1}")
```

```
accuracy: 0.4897959183673469
precision: 0.17721518987341772
```

```
recall: 0.5833333333333334
F1 score: 0.27184466019417475
```

1.0.16 T14. The majority rule is the accuracy if you use the most frequent class from the training set as the classification decision. Report the overall Accuracy. Then, report the Precision, Recall, and F score for attrition prediction using the majority rule baseline.

```
[ ]: test_size = y_test.shape[0]

most_appear_class = np.argmax(np.bincount(np.array(y_train, dtype=int)))

y_majority_result = np.array([most_appear_class] * test_size)

[ ]: accuracy, precision, recall, F1, fpr = evaluate(y_test, y_majority_result,
↪cls=1)

print(f"accuracy: {accuracy}")
print(f"precision: {precision}")
print(f"recall: {recall}")
print(f"F1 score: {F1}")
```

```
accuracy: 0.8367346938775511
precision: nan
recall: 0.0
F1 score: nan
```

```
<ipython-input-376-48a16e509fc3>:9: RuntimeWarning: invalid value encountered in
long_scalars
precision = np.sum(result[y_pred_arr == cls]) / np.sum(y_pred_arr == cls)
```

1.0.17 T15. Compare the two baselines with your Naive Bayes classifier.

- Naive Bayes Classifier which calculate probability from histogram has better accuracy than Random classifier baseline and Majority rule baseline but has poor recall compared to random classifier baseline
- Naive Bayes Classifier which calculate probability from gaussian distribution also has better accuracy than random and little difference from majority rule baseline but it has quite better precision than random baseline

1.0.18 T16. Use the following threshold values

\$ t = np.arange(-5,5,0.05) \$ ### find the best accuracy, and F score (and the corresponding thresholds)

```
[ ]: t = np.arange(-5, 5, 0.05)
```

```
[ ]: y_predicted, classifier_arr = model.predict(x_test, 0)

F1 = (0.0, 0.0)
accuracy = (0.0, 0.0)

for threshold in t:
    result = classifier_arr <= threshold

    y_predicted[result == True] = 1
    y_predicted[result == False] = 0

    y_pred_arr = np.array(y_predicted)
    y_true_arr = np.array(y_test)

    fpr = y_true_arr == y_pred_arr

    _accuracy, _, _, _F1, fpr = evaluate(y_test, y_predicted, cls=1)

    if accuracy[0] < _accuracy:
        accuracy = (_accuracy, threshold)

    if F1[0] < _F1:
        F1 = (_F1, threshold)

print(f"Best Accuracy: {accuracy[0]} at threshold {accuracy[1]}")
print(f"Best F1 Score: {F1[0]} at threshold {F1[1]}")
```

Best Accuracy: 0.8775510204081632 at threshold -1.7500000000000115

Best F1 Score: 0.5652173913043478 at threshold 0.5499999999999803

1.0.19 T17. Plot the RoC of your classifier.

```
[ ]: y_predicted, classifier_arr = model.predict(x_test, 0)

roc_curve_tpr = []
roc_curve_fpr = []

for threshold in t:
    result = classifier_arr <= threshold


    y_predicted[result == True] = 1
    y_predicted[result == False] = 0

    y_pred_arr = np.array(y_predicted)
    y_true_arr = np.array(y_test)

    _, _, tpr, _, fpr = evaluate(y_test, y_predicted, cls=1)
```

```
roc_curve_tpr.append(tpr)
roc_curve_fpr.append(fpr)
```

```
[ ]: plt.plot(roc_curve_fpr, roc_curve_tpr)
plt.show()
```



1.0.20 T18. Change the number of discretization bins to 5. What happens to the RoC curve? Which discretization is better? The number of discretization bins can be considered as a hyperparameter, and must be chosen by comparing the final performance.

```
[ ]: new_model = SimpleBayesClassifier(n_pos = np.sum(y_train == 1), n_neg = np.
    ↳sum(y_train == 0))
```

```
[ ]: new_model.fit_params(x_train, y_train, 5)
```

```
[ ]: new_y_predicted, new_classifier_arr = new_model.predict(x_test, 0)

new_roc_curve_tpr = []
new_roc_curve_fpr = []

for threshold in t:
```

```

result = new_classifier_arr <= threshold

new_y_predicted[result == True] = 1
new_y_predicted[result == False] = 0

y_pred_arr = np.array(new_y_predicted)
y_true_arr = np.array(y_test)

_, _, tpr, _, fpr = evaluate(y_test, new_y_predicted, cls=1)


new_roc_curve_tpr.append(tpr)
new_roc_curve_fpr.append(fpr)

```

```

[ ]: plt.plot(new_roc_curve_fpr, new_roc_curve_tpr)
plt.show()

```



We see that true positive rate and false positive rate is significantly increased from bin size 10 and overall trend look like that ROC curve has increased faster than ROC curve when bin size is 10. I think bin size 10 has better discretization than bin size 5 because at the same true positive rate value, bin size 10 tends to give less false positive rate than bin size 5 if we compare in ROC curve